

Importing the necessary libraries

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn import metrics
```

Data Collection and Processing

```
In [2]: car_dataset = pd.read_csv("C:\\Users\\Lenovo\\Downloads\\car data.csv")
```

```
In [3]: car_dataset.head()
```

```
Out[3]:
```

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	Transmissior
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manua
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manua
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manua
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manua
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manua

```
In [4]: car_dataset.shape
```

```
Out[4]: (301, 9)
```

```
In [5]: car_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        301 non-null   object
1   Year            301 non-null   int64
2   Selling_Price   301 non-null   float64
3   Present_Price   301 non-null   float64
4   Driven_kms      301 non-null   int64
5   Fuel_Type       301 non-null   object
6   Selling_type    301 non-null   object
7   Transmission    301 non-null   object
8   Owner           301 non-null   int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

```
In [6]: car_dataset.isnull().sum()
```

```
Out[6]: Car_Name      0
        Year         0
        Selling_Price 0
        Present_Price 0
        Driven_kms     0
        Fuel_Type      0
        Selling_type    0
        Transmission    0
        Owner          0
        dtype: int64
```

```
In [7]: #checking the distribution of categorical data
print(car_dataset.Fuel_Type.value_counts())
print(car_dataset.Selling_type.value_counts())
print(car_dataset.Transmission.value_counts())
```

```
Fuel_Type
Petrol    239
Diesel    60
CNG        2
Name: count, dtype: int64
Selling_type
Dealer    195
Individual 106
Name: count, dtype: int64
Transmission
Manual    261
Automatic  40
Name: count, dtype: int64
```

Encoding the categorical data

```
In [8]: car_dataset.replace({'Fuel_Type':{'Petrol':0, 'Diesel':1, 'CNG':2}}, inplace=True)
car_dataset.replace({'Selling_type':{'Dealer':0, 'Individual':1}}, inplace=True)
car_dataset.replace({'Fuel_Type':{'Manual':0, 'Automatic':1}}, inplace=True)
```

```
In [9]: car_dataset.head()
```

```
Out[9]:
```

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	Transmission
0	ritz	2014	3.35	5.59	27000	0	0	Manua
1	sx4	2013	4.75	9.54	43000	1	0	Manua
2	ciaz	2017	7.25	9.85	6900	0	0	Manua
3	wagon r	2011	2.85	4.15	5200	0	0	Manua
4	swift	2014	4.60	6.87	42450	1	0	Manua

Splitting the data and target

```
In [10]: X = car_dataset.drop(['Car_Name', 'Selling_Price', 'Transmission'], axis=1)
Y = car_dataset['Selling_Price']
```

```
In [11]: print(X)
```

	Year	Present_Price	Driven_kms	Fuel_Type	Selling_type	Owner
0	2014	5.59	27000	0	0	0
1	2013	9.54	43000	1	0	0
2	2017	9.85	6900	0	0	0
3	2011	4.15	5200	0	0	0
4	2014	6.87	42450	1	0	0
..
296	2016	11.60	33988	1	0	0
297	2015	5.90	60000	0	0	0
298	2009	11.00	87934	0	0	0
299	2017	12.50	9000	1	0	0
300	2016	5.90	5464	0	0	0

[301 rows x 6 columns]

Splitting training and test data

```
In [12]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.1, random_state=42)
```

Model Training

1. Linear Regression

```
In [13]: lin_reg_model = LinearRegression()
```

```
In [14]: lin_reg_model.fit(X_train, Y_train)
```

```
Out[14]: ▼ LinearRegression
LinearRegression()
```

Model Evaluation

```
In [15]: training_data_prediction = lin_reg_model.predict(X_train)
```

```
In [16]: error_score = metrics.r2_score(Y_train, training_data_prediction)
print("R Squared Error : ", error_score)
```

R Squared Error : 0.8716571597791489

Visualize the actual prices and predicted prices

```
In [17]: plt.scatter(Y_train, training_data_prediction)
plt.xlabel("Actual Price")
```

```
plt.ylabel("Predicted Price")  
plt.title("Actual Prices vs Predicted Prices")  
plt.show()
```




```
In [18]: test_data_prediction = lin_reg_model.predict(X_test)
```

```
In [19]: error_score = metrics.r2_score(Y_test, test_data_prediction)  
print("R Squared Error : ", error_score)
```

R Squared Error : 0.8340577535179488

```
In [20]: plt.scatter(Y_test, test_data_prediction)  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.title("Actual Prices vs Predicted Prices")  
plt.show()
```



2. Lasso Regression

```
In [21]: lass_reg_model = Lasso()
```

```
In [22]: lass_reg_model.fit(X_train,Y_train)
```

```
Out[22]: ▾ Lasso  
Lasso()
```

Model Evaluation


```
In [23]: training_data_prediction = lass_reg_model.predict(X_train)
```

```
In [24]: error_score = metrics.r2_score(Y_train, training_data_prediction)  
print("R Squared Error : ", error_score)
```

R Squared Error : 0.8424480718240741

Visualize the actual prices and predicted prices

```
In [25]: plt.scatter(Y_train, training_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Prices vs Predicted Prices")
plt.show()
```



```
In [26]: test_data_prediction = lass_reg_model.predict(X_test)
```

```
In [27]: error_score = metrics.r2_score(Y_test, test_data_prediction)
print("R Squared Error : ", error_score)
```

R Squared Error : 0.8709763132343402

```
In [28]: plt.scatter(Y_test, test_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Prices vs Predicted Prices")
plt.show()
```

